

SDE Quiz - 3

1. When comparing file systems and DBMS, which factors best explain *why DBMS consumes more storage space* for the same data size?
 - A. It pre-allocates schema and metadata
 - B. It stores data redundantly across multiple drives
 - C. It maintains statistical and optimization data
 - D. It compresses data to improve access time

Answer: A, C

2. A system designer decides to store medical records in plain files for space efficiency but finds query performance too low. Which trade-off is evident here?
 - A. Time efficiency traded for space efficiency
 - B. Consistency traded for concurrency
 - C. Reliability traded for heterogeneity
 - D. Persistence traded for scalability

Answer: A

3. Why can't file systems completely replace DBMS even though both use B+ trees?
 - A. File systems lack query optimization layers
 - B. DBMS supports multiple index types
 - C. File systems cannot handle transactional integrity
 - D. File systems store data only in row-oriented format

Answer: A, B, C

4. Which principle best summarizes why DBMS trades *space for time*?
 - A. It stores redundant metadata for faster query optimization
 - B. It reduces I/O by removing indexes
 - C. It allocates more disk space to reduce computation cost during execution
 - D. It compresses data to reduce query execution time

Answer: C

5. During query execution, which DBMS components collaborate to ensure both correctness and efficiency?
 - A. Query parser
 - B. Scheduler
 - C. Recovery manager

D. File management layer

Answer: A, B, C, D

6. The “Empty Table Phenomenon” demonstrates which architectural design trade-off?
 - A. Overhead from pre-allocating schema and statistics
 - B. Inefficient memory paging
 - C. Lack of metadata optimization
 - D. Poor concurrency control

Answer: A

7. In which scenarios would a **hybrid database** (SQL + NoSQL) be most beneficial?
 - A. E-commerce platforms with product metadata and real-time transactions
 - B. Static government record archives
 - C. Scientific data pipelines with schema evolution
 - D. Systems requiring both ACID and BASE guarantees

Answer: A, C, D

8. Which aspect of UML parallels the concept of “entities” in ER models?

- A. Classes
- B. Methods
- C. Namespaces
- D. Interfaces

Answer: A

9. How does transaction management differ between relational and NoSQL databases?
 - A. RDBMS supports multi-table atomicity
 - B. NoSQL ensures ACID compliance across collections
 - C. NoSQL typically supports document-level transactions
 - D. RDBMS transactions guarantee isolation and durability

Answer: A, C, D

10. Which scenario violates ACID but satisfies BASE?

- A. A social media post appearing before final consistency
- B. A banking transfer rolling back successfully

- C. A file write operation with journaling
- D. A transactional invoice update

Answer: A

11. Which database choices align best with these use cases?

- o (i) Real-time bidding system
- o (ii) Student information system
- o (iii) Log analytics platform

- A. Redis → (i)
- B. PostgreSQL → (ii)
- C. Hadoop-based Key-Value → (iii)
- D. XML DB → (i)

Answer: A, B, C

12. The “Curse of Dimensionality” primarily impacts which design dimensions?

- A. Index efficiency
- B. Query discrimination
- C. Schema normalization
- D. Data sparsity

Answer: A, B, D

13. What design patterns enable reliability in DBMS ecosystems?

- A. Recovery manager redundancy
- B. Predictable query optimization strategies
- C. Static schema enforcement
- D. Error detection through metadata auditing

Answer: A, B, D

14. Why are **B+ trees** preferred for file system and DBMS indexing over **B-trees**?

- A. Fixed node depth ensures predictable performance
- B. Data stored only in leaf nodes
- C. Higher branching factor for faster traversal
- D. Reduced best-case access time

Answer: A, B, C

15. Which type of data system is best suited for applications needing *flexibility* and *semi-structured* storage?

- A. Graph Database
- B. Relational DB
- C. Key-Value Store
- D. XML Database

Answer: D

16. Which statements about Key-Value stores are true in analytical workloads?

- A. Excellent for high-volume writes
- B. Suitable for single-record ACID transactions
- C. Efficient for distributed analytical reads
- D. Inefficient for value-based searches

Answer: A, C, D

17. Which database types best match their **strengths**?

- A. XML DB → Schema flexibility
- B. Key-Value → Write optimization
- C. RDBMS → Complex relationship handling
- D. Graph DB → Transactional reliability

Answer: A, B, C

18. What limitations make relational databases unsuitable for rapidly changing schemas?

- A. Schema must be predefined
- B. High restructuring overhead
- C. Rigid constraint enforcement
- D. Lack of support for polymorphic records

Answer: A, B, C, D

19. Which database selection would best fit a financial transaction system?

- A. ACID-compliant RDBMS
- B. BASE-oriented NoSQL
- C. Key-Value store
- D. Graph database

Answer: A

20. When comparing **XML** and **Key-Value** databases, which distinctions are correct?

- A. XML supports semi-structured hierarchy
- B. Key-Value stores are schema-less
- C. XML supports path-based navigation
- D. Key-Value supports relational joins

Answer: A, B, C

21. Which property of B+ trees ensures consistent access time?

- A. Data stored in all nodes
- B. Uniform leaf depth
- C. Variable branching factor
- D. Reduced pointer count

Answer: B

22. Which indicators help detect relational schema self-descriptiveness issues?

- A. Missing timeline/audit data
- B. Inability to distinguish derived tables
- C. Lack of metadata for indexes
- D. Redundant key duplication

Answer: A, B

23. In MapReduce, which factors influence whether a workload is write-heavy or read-heavy?

- A. Sorting requirements
- B. Data shuffling strategy
- C. Key distribution
- D. Output aggregation frequency

Answer: A, B, C, D

24. A log analytics system needs fast writes and occasional aggregate reads. Which DB type is best?

- A. XML Database
- B. Graph Database
- C. RDBMS
- D. Key-Value Store

Answer: D

25. Which key challenges cause scalability failure in RDBMS systems?

- A. Non-linear storage growth
- B. Metadata overhead
- C. Index explosion
- D. Weak consistency mechanisms

Answer: A, B, C

26. Which of these represent *trade-offs* in database technology selection?

- A. ACID vs. BASE
- B. Scalability vs. Consistency
- C. Space vs. Time
- D. Declarative vs. Imperative programming

Answer: A, B, C, D

27. In a DBMS, which component is responsible for interpreting user queries into execution plans?

- A. Query Parser and Processor
- B. File Manager
- C. Buffer Manager
- D. Recovery Manager

Answer: A

28. You create hundreds of empty tables for a multi-tenant prototype and storage spikes.

Which design choices explain the spike and help avoid it?

- A. Each empty table preallocates index structures and stats.
- B. Use a shared table with tenant_id column instead of many tables.
- C. Turn off query optimization statistics globally.
- D. Use partitioning or a single JSON/document column per tenant.

Answer: A, B, D

29. You must design a system that handles banking transactions (strict consistency), user profiles with dynamic fields, and high-throughput event logs. Which hybrid combination is most appropriate?

- A. RDBMS (ACID) for transaction processing
- B. Document store (schema-flexible) for user profiles
- C. Key-value store for event logs
- D. Graph DB for event logs

Answer: A, B, C

30. Which locking strategy in an XML DB reduces contention for updates to deep nodes while preserving parent-child integrity?

- A. Path locking entire root-to-leaf path for every update
- B. Node-level locking with subtree extension for parents
- C. Global database lock during each update
- D. No locking, rely on optimistic concurrency without checks

Answer: B

31. You need to design a session store with very fast reads and occasional bulk analytics on sessions. Which design decisions balance performance and analytics needs?

- A. Keep sessions in an in-memory key-value store with TTLs.
- B. Periodically snapshot sessions to a distributed key-value store with time-partitioning for analytics.
- C. Always store sessions directly in RDBMS to simplify analytics.
- D. Emit session events to a log system (MapReduce/Spark) for batch aggregation.

Answer: A, B, D

32. You must design a remote island data center (limited power, intermittent connectivity, hurricane risk) for a mixed workload. Which single strategy best mitigates both availability and data safety risks?

- A. Use a hybrid cloud architecture with local DB for fast ops and remote replicas when connectivity allows.
- B. Centralize everything in a single powerful server on the island.
- C. Use in-memory databases only to save power.
- D. Rely solely on NoSQL eventual-consistency stores to reduce replication overhead.

Answer: A

33. A database optimized for *read-heavy workloads* would likely adopt which design decisions?

- A. Denormalized tables
- B. Column-oriented storage
- C. Frequent index rebuilding
- D. Precomputed materialized views

Answer: A, B, D

34. Which database property most negatively impacts write performance in a traditional RDBMS?

- A. Normalization

- B. Lack of indexing
- C. ACID transaction enforcement
- D. Row-oriented storage

Answer: C

35. You're designing a database for IoT sensor logs with **billions of daily inserts** but infrequent reads. Which architectural patterns fit best?

- A. Append-only log architecture
- B. In-memory cache for recent writes
- C. Extensive secondary indexing
- D. NoSQL key-value model

Answer: A, B, D

36. A “both read and write not friendly” system is often a result of:

- A. Poor schema design
- B. Over-normalization
- C. Over-indexing with high maintenance overhead
- D. Balanced indexing and normalization

Answer: A, B, C

37. You're asked to justify constraints for a *write-heavy, read-light* system. Which constraints are likely to be **relaxed or modified**?

- A. Immediate foreign key enforcement
- B. Strict ACID durability
- C. Index creation on non-primary attributes
- D. Eventual consistency adoption

Answer: A, B, C, D

38. Which real-world system is **read-heavy but write-light**?

- A. E-commerce order system
- B. Social media feed viewer
- C. Stock trading order book
- D. IoT sensor data stream

Answer: B

39. To achieve *write efficiency* while maintaining acceptable query performance, a schema designer might:

- A. Use partitioning or sharding
- B. Avoid multi-table joins
- C. Normalize all attributes into separate tables
- D. Use batch writes instead of individual commits

Answer: A, B, D

40. Match these real-world systems with their likely database category:

- A. Banking core system → Both read/write friendly
- B. Search engine index → Read friendly, write unfriendly
- C. Log collection system → Write friendly, read unfriendly
- D. Historical archive → Both read/write unfriendly

Answer: A, B, C, D