

**1. A banking system needs to ensure no transaction loss or inconsistency even in case of power failure. Which property of DBMS is most relevant?**

- A. Concurrency
- B. Durability
- C. Non-redundancy
- D. Efficiency

**Answer:** B — Durability ensures transaction results persist despite failures.

**2. A distributed database spanning multiple continents suffers from network partitioning. Which CAP properties can realistically be preserved simultaneously?**

- A. Consistency and Availability
- B. Availability and Partition Tolerance
- C. Consistency and Partition Tolerance
- D. All three

**Answer:** B and C — CAP theorem forbids all three; systems must trade off between C-A-P.

**3. If a company must scale rapidly with heterogeneous data formats (JSON, CSV, sensor data), which database type suits best?**

- A. Relational DBMS
- B. Graph Database
- C. Document Store
- D. Column Store

**Answer:** C — Document DBs handle heterogeneous, schema-flexible data effectively.

**4. In an e-commerce platform, which anomalies can occur if normalization is not properly applied?**

- A. Redundant data
- B. Data inconsistency
- C. Faster query performance
- D. Deletion anomaly

**Answer:** A, B, D — Denormalization causes redundancy, inconsistency, and anomalies.

**5. A transaction scheduler enforces 2-phase locking. What is the main trade-off?**

- A. Higher throughput, weaker isolation
- B. Lower concurrency, stronger isolation
- C. No deadlocks
- D. Better recovery control

**Answer:** B — 2PL enforces serializability but limits concurrency.

**6. Which modeling method would best represent a social network with relationships and multiple entity types?**

- A. E-R Model
- B. UML Class Diagram
- C. Graph Model
- D. Relational Schema

**Answer:** C — Graph models handle rich relationship semantics.

**7. You are designing a high-read, low-write analytical system. Which data organization strategies help performance?**

- A. Column-oriented storage
- B. Row-oriented storage
- C. Data compression (e.g., run-length encoding)
- D. Schema normalization

**Answer:** A, C — Columnar storage with compression optimizes read-heavy workloads.

**8. A company needs real-time analytics from distributed logs. Which system type fits best?**

- A. Document Database
- B. Graph Database
- C. Column Store
- D. Distributed Key-Value Store

**Answer:** D — Key-value stores scale horizontally for fast distributed reads/writes.

**9. Which relational operation would you use to retrieve all employees working in both HR and Finance departments?**

- A. Projection
- B. Union
- C. Intersection
- D. Join

**Answer:** C — Intersection retrieves common tuples across relations.

**10. A database uses timestamp ordering for concurrency control. What scenario can still cause rollback?**

- A. Conflicting writes with older timestamps
- B. Serial execution

- C. Read-only transactions
- D. Use of strict two-phase locking

**Answer:** A — Timestamp conflicts may cause transaction aborts/rollbacks.

**11. You must migrate from a relational schema to a graph database. What adjustments are likely needed?**

- A. Introduce node/edge identifiers
- B. Flatten recursive relationships
- C. Move from fixed schema to property-based model
- D. Use tuple-based joins instead of traversal

**Answer:** A, C — Graph DBs rely on nodes/edges with flexible properties.

**12. In a distributed database, replication transparency ensures:**

- A. Each copy must be identical
- B. User queries behave the same regardless of replica location
- C. Updates propagate only manually
- D. Network failures are hidden

**Answer:** B — Replication transparency hides replica management from users.

**13. When is a weakly consistent system acceptable?**

- A. In financial ledgers
- B. In live stock trading
- C. In social media feeds
- D. In bank transfer logs

**Answer:** C — Eventual consistency suffices for social applications.

**14. A sensor data platform stores billions of readings with variable schema. Which features are beneficial?**

- A. Schema independence
- B. Versioned data
- C. Strict normalization
- D. Scalability

**Answer:** A, D — Schema flexibility and scalability matter most.

---

**15. Which join will preserve all tuples from the left relation even if no match exists?**

- A. Inner Join
- B. Natural Join
- C. Left Outer Join
- D. Semi Join

**Answer:** C — Left outer join includes unmatched left tuples.

**16. You observe semantic overloading in a relational schema. What is the probable issue?**

- A. Too many tables
- B. One attribute used to store multiple meanings
- C. High redundancy
- D. Weak relationships

**Answer:** B — Semantic overloading means attributes lose precise meaning.

**17. For an XML-based data interchange between systems, which standards are most critical?**

- A. XQuery
- B. XML Schema
- C. XSLT
- D. XPath

**Answer:** A, B, C — Schema defines structure; XQuery/XSLT handle querying and transformation.

**18. In a MapReduce job, which phase is responsible for combining values with identical keys?**

- A. Map
- B. Shuffle
- C. Reduce
- D. Split

**Answer:** C — Reduce phase aggregates intermediate key-value pairs.

**19. A graph DB model supports weighted edges and multiple edge types. Which model best fits?**

- A. Property Graph
- B. Simple Directed Graph
- C. Hypergraph
- D. Nested Graph

**Answer:** A, D — Property graphs support metadata; nested graphs handle hierarchy.

**20. Which factor primarily drives the need for distributed databases?**

- A. Query simplicity
- B. Redundancy reduction
- C. Scalability and fault tolerance
- D. Security

**Answer:** C — Distributed DBs are designed for scalable and fault-tolerant architectures.

**21. A hospital database frequently updates patient records but rarely deletes them. Which storage strategy will minimize write amplification and improve access?**

- A. Row-oriented storage
- B. Column-oriented storage
- C. Graph storage
- D. Document storage

**Answer:** A — Row stores handle frequent updates efficiently since tuples are co-located.

**22. A relational schema is normalized into 3NF. However, query performance drops. What are valid reasons?**

- A. Excessive joins due to decomposed tables
- B. Larger indexes
- C. Reduced data redundancy
- D. Poor caching locality

**Answer:** A, D — More joins and fragmented data reduce performance despite theoretical soundness.

**23. In a distributed system using eventual consistency, which statement is TRUE?**

- A. All nodes always see the same data instantly
- B. Some reads may return stale data temporarily
- C. Writes are serialized globally
- D. Data replication is synchronous

**Answer:** B — Eventual consistency allows temporary divergence.

**24. If a database designer mistakenly stores “TotalSalary” as an attribute derived from other columns, what anomaly may arise?**

- A. Update anomaly
- B. Insertion anomaly

- C. Deletion anomaly
- D. Redundancy anomaly

**Answer:** A — Derived fields must be recomputed; inconsistent updates cause anomalies.

**25. A graph traversal query performs poorly. What optimizations could help?**

- A. Use adjacency lists instead of matrices
- B. Store frequently accessed edges in cache
- C. Introduce redundant edge labels
- D. Denormalize graph structure

**Answer:** A, B — Adjacency lists and caching reduce traversal cost.

**26. Which feature of DBMS ensures rollback to last consistent state after a crash?**

- A. Scheduler
- B. Recovery Manager
- C. Buffer Manager
- D. Query Parser

**Answer:** B — Recovery Manager handles log-based rollback and restart.

**27. A system requires high concurrency but is tolerant to temporary inconsistencies. Which isolation level fits?**

- A. Serializable
- B. Repeatable Read
- C. Read Committed
- D. Read Uncommitted

**Answer:** C — Balances concurrency with moderate isolation.

**28. A column store compresses data using multiple encodings. Which properties influence encoding choice?**

- A. Value distribution
- B. Column cardinality
- C. Query workload
- D. Tuple size

**Answer:** A, B, C — Encoding must suit data patterns and access type.

**29. Which XML concept is loosely analogous to SQL's SELECT-FROM-WHERE clause?**

- A. XSLT
- B. XPath
- C. FLWOR Expression

## D. XML Schema

**Answer:** C — FLWOR defines query expressions like SQL.

**30. In a NoSQL key-value store, you observe skewed load on some nodes. What techniques can mitigate this?**

- A. Consistent hashing
- B. Sharding by key prefix
- C. Randomized partitioning
- D. Static partitioning

**Answer:** A, C — Distributes keys more evenly across nodes.

**31. Which failure scenario primarily tests network partition tolerance?**

- A. Server crash
- B. Message delay across data centers
- C. Schema migration
- D. Disk failure

**Answer:** B — Network partitions test CAP trade-offs between consistency and availability.

**32. For a geo-distributed social media platform, which strategies enhance both latency and availability?**

- A. Data replication near users
- B. Asynchronous updates
- C. Strong global consistency
- D. Write-ahead logging

**Answer:** A, B — Replication and async writes improve performance while tolerating delays.

**33. In graph databases, what happens if you replace a “friend” edge with a “follower” edge but forget to update the reverse direction?**

- A. Data redundancy
- B. Loss of referential integrity
- C. Deadlock
- D. Serialization failure

**Answer:** B — Relationships lose bidirectional integrity.

**34. Which challenges are more pronounced in NoSQL systems compared to relational databases?**

- A. Lack of strict schema
- B. Weak consistency
- C. Complex joins
- D. Atomicity loss

**Answer:** A, B, D — NoSQL trades strictness for flexibility and scalability.

**35. Which graph traversal ensures every node is visited exactly once before returning to the start?**

- A. Eulerian Cycle
- B. Hamiltonian Cycle
- C. BFS
- D. Depth-limited Search

**Answer:** B — Hamiltonian cycles visit all vertices exactly once.

**36. Which mechanisms can prevent lost updates in concurrent transactions?**

- A. Two-phase locking
- B. Timestamp ordering
- C. Deferred writes
- D. Shadow paging

**Answer:** A, B — Both enforce serializable scheduling.

**37. If an XML document needs to be transformed into HTML, which standard is most applicable?**

- A. XPath
- B. XSLT
- C. XQuery
- D. XML Schema

**Answer:** B — XSLT transforms XML into other formats like HTML.

**38. A columnar store using run-length encoding would perform best when:**

- A. Column values are unique
- B. Column values are highly repetitive
- C. Query involves random writes

- D. Tuple reconstruction is frequent

**Answer:** B — Run-length encoding compresses repeated values efficiently.

**39. In distributed databases, which failures can cause partial transaction commit?**

- A. Message failure
- B. Server crash
- C. Network partition
- D. Disk failure

**Answer:** A, B, C — Communication and node failures disrupt distributed commits.

**40. Which consistency model ensures that once you've read your own write, you'll never see an older version again?**

- A. Monotonic Writes
- B. Monotonic Reads
- C. Read Your Writes
- D. Causal Consistency

**Answer:** B — Monotonic reads guarantee non-regression in read versions.

**41. A document database exposes a REST API with CRUD operations. What advantage does REST's stateless nature provide?**

- A. Faster joins
- B. Easier load balancing
- C. Guaranteed consistency
- D. Fewer network calls

**Answer:** B — Statelessness enables independent request handling and load distribution.

**42. For a multi-tenant SaaS platform, what design patterns improve scalability?**

- A. Database sharding
- B. Denormalization
- C. Data federation
- D. Centralized transaction manager

**Answer:** A, B, C — Distributing data and simplifying queries boosts scalability.

**43. In a system emphasizing availability over consistency, what issue might arise after network partition recovery?**

- A. Lost updates

- B. Data divergence
- C. Reduced throughput
- D. Schema lock

**Answer:** B — Divergent replicas may exist temporarily or permanently.

**44. Which concepts from UML are directly mappable to ER modeling?**

- A. Class ↔ Entity
- B. Attribute ↔ Column
- C. Association ↔ Relationship
- D. Method ↔ Constraint

**Answer:** A, B, C — UML classes/attributes/associations align with ER entities.

**45. What is the primary benefit of using a buffer manager in DBMS?**

- A. Query parsing
- B. Minimizing disk I/O
- C. Enforcing foreign keys
- D. Handling deadlocks

**Answer:** B — Buffer managers reduce disk access by caching frequently used pages.

**46. To guarantee atomic commit in distributed transactions, which protocols are essential?**

- A. Two-Phase Commit
- B. Gossip Protocol
- C. Write-Ahead Logging
- D. Consensus Algorithms (e.g., Paxos)

**Answer:** A, D — Both ensure agreement on commit across nodes.

**47. Why might homogeneous structure in relational tables be a limitation?**

- A. Difficult to normalize
- B. Restricts hierarchical data representation
- C. Increases redundancy
- D. Slows indexing

**Answer:** B — Fixed schema restricts nested or irregular data modeling.

**48. A property graph can represent multiple edge types such as “`works_with`” and “`reports_to`.” What does this achieve?**

- A. Schema rigidity
- B. Semantic richness
- C. Easier traversal
- D. Query optimization

**Answer:** B, C — Multiple relation types enrich semantics and improve navigation.

**49. During map phase of MapReduce, if keys are not properly partitioned, what could happen?**

- A. Incorrect aggregation
- B. Shuffle phase failure
- C. Skewed reducer load
- D. Network congestion reduction

**Answer:** C — Poor key partitioning creates load imbalance.

**50. A distributed document database requires strong consistency for financial data and high availability for logs. How should the architecture handle this?**

- A. Use CAP trade-offs per data type
- B. Replicate financial data synchronously
- C. Use eventual consistency for logs
- D. Force global ACID compliance

**Answer:** A, B, C — Hybrid consistency models satisfy different workloads.